



# 046273 Distributed Functional Programming


Lecture 1: Welcome to Erlang

*Hovav Gazit, Winter 2017*

*Based on slides of Prof' Idit Keidar*

# Today's Agenda

- A little about this course
  - Motivation
  - Formalities
- Getting started with Erlang
  - What is Erlang
  - "Hello World" in Erlang
  - The Erlang shell
  - Types: numbers, atoms, booleans, tuples, lists
  - Variables & pattern matching
  - Simple functions & recursion



# What is This Course About?

1. The Erlang approach to concurrency
    - Functional programming
    - Many lightweight processes
    - Asynchronous message-passing
      - No shared memory
    - Distributed systems
    - Fault-tolerance, robustness
  2. The principles behind these features
- Let's start with a little motivation ...

# The (Computing) World is Concurrent

- The multi-core revolution  $\Rightarrow$  computing with multiple concurrent processes/threads - **parallel computing**
- Wide dispersion, scalability, fault tolerance  $\Rightarrow$  **distributed computing**
- Inter-operability, interfacing with internal and external systems  $\Rightarrow$  **middleware** for distributed communication & coordination

# The World is Concurrent

"The world is parallel. If we want to write programs that behave as other objects behave in the real world, then these programs will have a concurrent structure."

**-- Joe Armstrong**

*To be continued ...*



# Dealing with Concurrency


- Many tools, languages, paradigms,...
- For parallel computing – most commonly using shared memory
  - pthreads (see SOS), transactional memory, ...
  - error prone, hard to get right
- For distributed computing – usually using message-passing
  - sockets, message queues, JMS
  - inter-operability using middleware
  - reliability always a challenge

# Dealing with Concurrency Take II

“  
...

Use a language that was designed for writing concurrent applications, and development becomes a lot easier. Erlang programs model how we think and interact.”

-- **Joe Armstrong**, one of the creators of Erlang



# Course Staff

- Lecturer (& TA): Hovav Gazit
  - [hovav@technion.ac.il](mailto:hovav@technion.ac.il)
  - Room 1009 (Inside CGM lab, Room 1011)
  - Appointment by email
  - Phone: 829-5741
- Homework Checker: Itay Tsabary
- Academic Supervisor: Prof' Idit Keidar
  - [idish@ee.technion.ac.il](mailto:idish@ee.technion.ac.il)
- Course web site: Moodle



# Requirements and Grading

- Erlang programming assignments 25%
  - 3 - 4 Wet Exercises
- Exam 75%
  - Allowed material: lecture notes, and your handwritten notes *only*
  - Moed A: 2/2 (Thursday)
  - Moed B: 1/3 (Wednesday)
- In order to pass the course you need to pass the EXAM



# Prerequisite Background



- Some programming experience
  - e.g., in C/C++/Java
  - be comfortable reading reference **manuals**
- Some background in algorithms and data structures highly recommended
  - feel comfortable with **recursion**
- **MAMAT & Operating Systems (In parallel)**
- Have you heard about other approaches to **concurrency**?
  - OS - pthreads

# Resources

- Textbooks:
  - **Programming Erlang: Software for a Concurrent World**, Joe Armstrong
  - I use code examples from the book <http://pragmaticprogrammer.com/titles/jaerlang/code.html>.
  - **Erlang Programming - A Concurrent Approach to Software Development**, Francesco Cesarini, Simon Thompson  
O'Reilly Media 2009
- Erlang official web site <http://www.erlang.org/>
  - Download Erlang, Reference Manual
- Tutorial (I use pictures, examples): <http://learnyouosomeerlang.com/contents>

Programming  
Erlang Software for a  
Concurrent World



Joe Armstrong



# What is Erlang?

A programming language for concurrent software



Agner Krarup Erlang (1878-1929)  
Danish mathematician invented  
fields of traffic engineering and  
queuing theory



Ericsson  
Language

# Erlang Language Characteristics

- Functional language
  - High-level, compact **declarative** programming
  - Supports high-order functions
  - No for, while loops – use recursion instead
  - Immutable state – use tail recursion instead
- Garbage-collected
  - No explicit memory management
- Runs on virtual machine
  - Source compiled into bytecode

# Erlang Concurrency Features

- Scalable, lightweight concurrency
- Based on message passing
  - no shared memory
- Soft real-time
- Robust, fault-tolerant
- Supports distributed computations
- Open, integrates with other platforms



# Erlang is Best For

- “If your target system is a
    - high-level, concurrent, robust, soft real-time system
    - that will scale in line with demand,
    - make full use of multi-core processors, and
    - integrate with components written in other languages
  - Erlang should be your choice.”
- **Cesarini and Thompson**

# Erlang is Used By

- [Amazon](#) SimpleDB, DB part of EC2
- [Yahoo!](#) Delicious social bookmarking
  - > 5 million users and 150 million bookmarked URLs
- [Facebook](#) chat service backend,
  - > 100 million active users
- [T-Mobile](#) SMS & authentication
- [Motorola](#) call processing in public-safety industry
- [Ericsson](#) GPRS and 3G mobile networks worldwide
- [WhatsApp](#) “How do you support 450 million users with only 32 engineers ?”
- Many open source projects



# 5 Reasons to Learn Erlang

## by Joe Armstrong

1. You want to write programs that run faster when you run them on a **multi-core** computer.
2. You want to write **fault-tolerant** applications that can be modified without taking them out of service.
3. You've heard about "**functional programming**" and you're wondering whether the techniques really work.
4. You want to use a language that has been battle tested in **real large-scale industrial products** that has great libraries and an active user community.
5. You don't want to wear your fingers out by typing lots of **lines of code**.

6. It is fun!

# Hello World in Erlang

```
-module(world).  
-export([hello/0]).
```

Module name

“Public” function

Comment

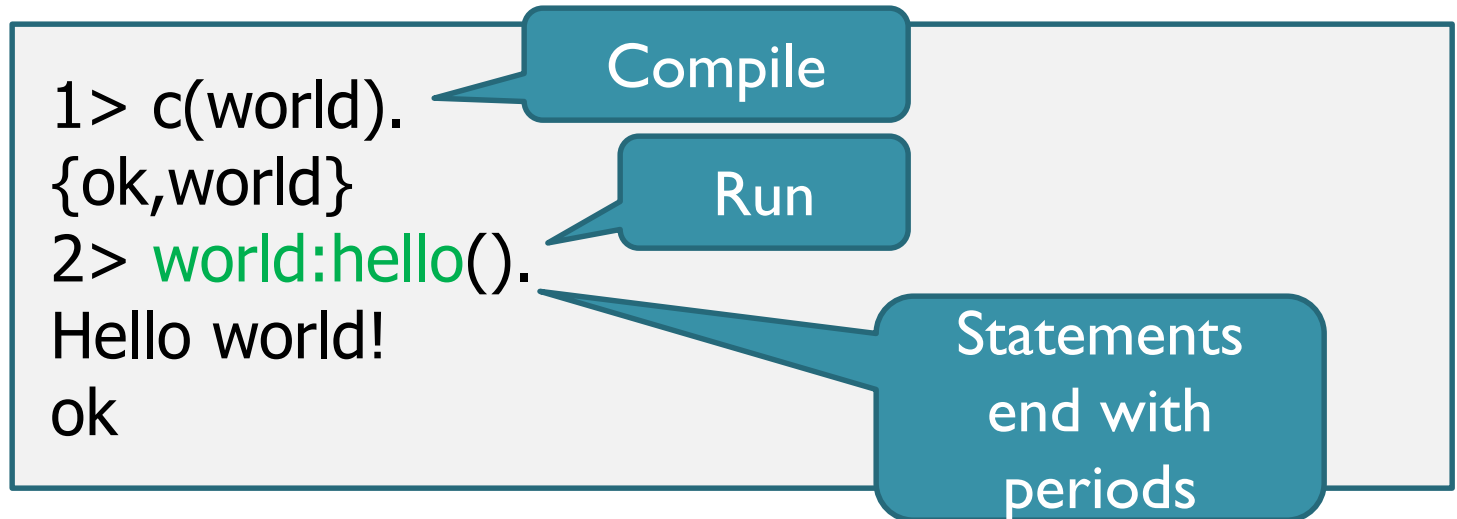
```
% Say hello to the world  
hello() -> io:format("Hello world!~n").
```

- Erlang code consists of modules
- This module's name is 'world'
  - must be stored in file 'world.erl'

Statements  
end with  
periods

# The Erlang Shell

- See [erlang.org](http://erlang.org) for installing and running
- Let's compile and run our new program:



We'll now learn about Erlang types and variables, and evaluate some expressions in the shell



# Erlang Basic Types, Variables, Pattern Matching

Chapter 2 of Francesco Cesarini

Chapter 2 of Armstrong

# Numbers

- **Integers**, can be *any* size
  - e.g., 18446744073709551616 ( $2^{64}$ ), -100
- **Floats**: double-precision
  - e.g., 6.148914691236517e18
- Evaluating arithmetic expressions in shell:

1> -234 + 1.

-233

2> 5-

2> 4.

1

3> 2#111.

7

4> (12+3)/5 + 7\*2.

17.0

Erlang evaluates the expression  
only when you enter a period

111 in base 2

The result is a float

# Arithmetic Expressions

<i>Op</i>	<i>Description</i>	<i>Argument Type</i>	<i>Priority</i>
$+ X$	$+ X$	Number	1
$- X$	$- X$	Number	1
$X * Y$	$X * Y$	Number	2
$X / Y$	$X / Y$ (floating-point division)	Number	2
$bnot X$	Bitwise not of $X$	Integer	2
$X div Y$	Integer division of $X$ and $Y$	Integer	2
$X rem Y$	Integer remainder of $X$ divided by $Y$	Integer	2
$X band Y$	Bitwise and of $X$ and $Y$	Integer	2
$X + Y$	$X + Y$	Number	3
$X - Y$	$X - Y$	Number	3
$X bor Y$	Bitwise or of $X$ and $Y$	Integer	3
$X bxor Y$	Bitwise xor of $X$ and $Y$	Integer	3
$X bsl N$	Arithmetic bitshift left of $X$ by $N$ bits	Integer	3
$X bsr N$	Bitshift right of $X$ by $N$ bits	Integer	3

Operators with equal priorities are left associative - evaluated from left to right.

# Variables Don't Vary

- **Variables** begin with an **upper case letter** (or **\_**), include **letters**, **digits**, **\_**
- Erlang is a single-assignment language

```
1> X = 500.
```

```
500
```

```
2> X.
```

```
500
```

```
3> Y.
```


**\* 1: variable 'Y' is unbound**

```
4> Y = X*X.
```

```
250000
```

```
5> X = 400.
```

**\*\* exception error: no match of right hand side value 400**



# Think of = in Algebraic Equations

- $X^2 = X + 2$ 
  - This must be the same  $X$  on both sides
- $X = X + 1$ 
  - This is impossible!
- $\begin{cases} X = 500 \\ Y = X^2 \end{cases}$ 
  - $Y$  has to stay the square of  $X$



# Pattern Matching

- = is *not* an assignment operator
- It's a **pattern matching** operator:

**Pattern = Term**

1> X = 5.

5

2> Y = 2 + 3.

5

3> Y = X.

5

4> X = Y + 1.

**\*\* exception error: no match of right hand side value 6**

5> f(X).

ok

6> X = Y + 1.


6

Can include unbound variables - will become bound

No unbound variables

Forget X

# Atoms



- **Atoms** are non-numerical constants
  - Like enums, or #defined symbols in C
- Start with **lower case letter**, include **letters, digits, @, \_, .**
  - examples: january,  
start\_with\_lower\_case,  
erlang@ee.technion.ac.il
- Or anything inside quotes
  - examples: 'January' 'spaces allowed'

# Booleans



- The atoms **true** and **false**.
- Boolean expressions:

1> 1 == 2.

false

2> 1 =< 2.

true

3> 1.0 == 1.

true

4> 1.0 == 1.

false

5> a < z.

true

6> (1<3) and (3 <4).

true

== means equals  
:= means is identical to

Atoms compared in  
lexicographical order


# Boolean Expressions and Term Comparisons

- **not** B1: Logical not
- B1 **and** B2: Logical and
- B1 **or** B2: Logical or
- B1 **xor** B2: Logical xor

<i>Operator</i>	<i>Meaning</i>
$X > Y$	X is greater than Y.
$X < Y$	X is less than Y.
$X \leq Y$	X is equal to or less than Y.
$X \geq Y$	X is greater than or equal to Y.
$X == Y$	X is equal to Y.
$X \neq Y$	X is not equal to Y.
$X ::= Y$	X is identical to Y.
$X \neq Y$	X is not identical to Y.

# Tuples

- Like struct, but anonymous



```
1> F = {firstName, joe}.
{firstName,joe}
2> L = {lastName, armstrong}.
{lastName,armstrong}
3> P = {person, F, L}.
{person,{firstName,joe},{lastName,armstrong}}
4> tuple_size(P).
3
5 > element(2, P).
{firstName,joe}
6> setelement(2, P, {firstname, joey}).
{person,{firstname,joey},{lastName,armstrong}}
```

Built-in functions (BIFs)  
manipulating tuples

# Pattern Matching with Tuples

- We can use pattern matching to extract values from tuples

```
1> Point = {point, 7.3, 45.2}.
```

```
{point, 7.3, 45.2}.
```

```
2> {point, X, Y} = Point.
```

```
{point,7.3,45.2}
```

```
3> X.
```

```
7.3
```

```
4> {point, C, C} = Point.
```

```
** exception error: no match of right hand side value
```

```
{point,7.3,45.2}
```

# Anonymous Variables in Patterns

- `_` is an **anonymous** variable
- Unlike regular variables, several occurrences of `_` in the same pattern don't have to bind to the same value

```
1> Guy = {person, {name, {first, joe}, {last, armstrong}}, {foot, 42}}.  
{person, {name, {first, joe}, {last, armstrong}}, {foot, 42}}  
2> {_, {_, {_, Who}, _}, _} = Guy.  
{person, {name, {first, joe}, {last, armstrong}}, {foot, 42}}  
3> Who.  
joe
```


# Uses of Pattern Matching

- Assigning values to variables
  - see “Variables ...” slide
- Extracting values from data structures
  - see last two slides
- Controlling execution flow
  - later – see “Functions with Multiple Clauses”



# Lists

- Lists look a lot like tuples
  - but are **processed** differently
- Example: `[a,b,c,d,177]`
  - `a` is the **head** of the list,
  - `[b,c,d,177]` is its **tail**
- The tail of a well-formed list
  - is a list
- `[]` is the **empty** list
- List **constructor** (**cons** for short):
  - `[H|T]` constructs a list with head `H` and tail `T`
  - e.g.,  $[1,2,3] = [1|[2,3]] = [1,2|[3]] = [1,2|[3|[]]]$



# List Concatenation, Subtraction

```
1> Part1 = [1,2,3],  
1> Part2 = [4,5,6],  
1> List = Part1 ++ Part2.
```

```
[1,2,3,4,5,6]
```

```
2> [a, a, b, b, c, c] -- [a, b, c].
```

```
[a,b,c]
```

```
3> [a, a, b, b, c, c] -- [a, b, c, b].
```

can you guess?


Multiple expressions  
separated by commas

List concatenation

List subtraction

- ++ and -- are right associative
- Note: these operators are inefficient
  - cons is more efficient
  - no need to worry about it in this course 😊

# Strings are Lists



```
1> Name = "Hello".
"Hello"
2> [1,2,3].
[1,2,3]
3> [83,117,114,112,114,105,115,101].
"Surprise"
4> [1,83,117,114,112,114,105,115,101].
[1,83,117,114,112,114,105,115,101].
5> I = $s.
115
6> [I-32,$u,$r,$p,$r,$i,$s,$e].
"Surprise"
7> "A long string"
7> " is split".
"A long string is split"
```

# Pattern Matching **pop?** **quiz**

Pattern	Term	Result?
{X,abc}	= {123,abc}	Succeeds, $X \rightarrow 123$
{X,Y,Z}	= {222,def,"cat"}	
{X,Y}	= {333,ghi,"cat"}	
X	= true	
{X,Y,X}	= {{abc,12},42,{abc,12}}	
{X,Y,X}	= {{abc,12},42,true}	
[H T]	= [1,2,3,4,5]	
[H,T]	= [1,2,3,4,5]	
[H T]	= "cat"	
[A,B,C T]	= [a,b,c,d,e,f]	

Call `f()` after each line, so all variables are unbound.



# Erlang Functions

Chapters 2, 3 of Francesco Cesarini

Chapter 3 of Armstrong

# Defining Functions



- Defined in **modules** (not in shell)
- Uniquely defined by
  - module name, function name, and **arity**
  - **arity** = number of arguments
- Argument, return types not specified
  - dynamic typing
  - checked at run-time

# Hello Function Example

- Let's recall our example:

`hello()` -> `io:format("Hello world!~n")`.

- This function is defined using 1 clause
- The head includes:
  - `'hello'` – the function name (an Atom)
  - `'()'` –there are no arguments, the *arity* is 0
- After -> comes the body
  - it invokes the function `'format/1'` (arity 1) from the module `'io'`

# Function Example 2

```
-module(demo).  
-export([print_square/1]).  
print_square(X) ->  
    Sq = X * X,  
    io:format("The square of ~p is ~p.~n", [X, Sq] ).
```

The arity of  
print\_square is 1

Expressions are  
separated by commas

The 2<sup>nd</sup> argument  
of format/2 is a list

Last expression  
ends with period

```
1> c(demo).  
{ok,demo}  
2> demo:print_square(7).  
The square of 7 is 49.  
ok  
3> demo:print_square(1.5).  
The square of 1.5 is 2.25.
```



# Functions with Multiple Clauses

```
-module(geometry).
```

```
-export([area/1]).
```

The arity of area is 1  
as it takes a single tuple

```
area({rectangle, Width, Height}) ->
```

```
    Width * Height;
```

Pattern matching : find  
first matching clause

```
area({square, Edge}) ->
```

```
    Edge * Edge;
```

Clauses separated by  
semicolons ;

```
area({circle, R}) ->
```

```
    math:pi() * R * R.
```

Last clause ends with period

```
1> geometry:area({circle, 1.4}).
```

```
6.157521601035994
```

```
2> geometry:area(1.4).
```

```
guess what?
```

# How Would We Write This in C?

```
enum ShapeType { Rectangle, Circle, Square };  
struct Shape {  
    enum ShapeType kind;  
    union {  
        struct { int width, height; } rectangleData;  
        struct { int radius; } circleData;  
        struct { int side;} squareData;  
    } shapeData;  
};  
double area(struct Shape* s) {  
    if( s->kind == Rectangle ) {  
        int width, ht;  
        width = s->shapeData.rectangleData.width;  
        ht = s->shapeData.rectangleData.ht;  
        return width * ht;  
    } else if ( s->kind == Circle ) {  
        ...  
    }  
}
```

# Recursion



Start off by defining  
base case(s)

`factorial(0) -> 1;`

`factorial(N) -> N * factorial(N - 1).`

Other cases  
should converge  
to base case(s)

# A Run of Factorial

`factorial(0) -> 1;`

`factorial(N) -> N * factorial(N - 1).`

$\text{factorial}(3) \Rightarrow 3 * \text{factorial}(3-1)$

$\text{factorial}(2) \Rightarrow 2 * \text{factorial}(2-1)$

$\text{factorial}(1) \Rightarrow 1 * \text{factorial}(1-1)$

$\text{factorial}(0) \Rightarrow 1$

$\Rightarrow 1 * 1 = 1$

$\Rightarrow 2 * 1 = 2$

$\Rightarrow 3 * 2 = 6$

The first clause matches

# Recursion Over Lists

`sum([]) -> 0;`

Base case

`sum([Head | Tail]) -> Head + sum(Tail).`

What is the arity of sum?

$\text{sum}([1,2,3]) \Rightarrow 1 + \text{sum}([2, 3])$

$\text{sum}([2,3]) \Rightarrow 2 + \text{sum}([3])$

$\text{sum}([3]) \Rightarrow 3 + \text{sum}([])$

$\text{sum}([]) \Rightarrow 0$

$3 + 0 \Rightarrow 3$

$2 + 3 \Rightarrow 5$

$1 + 5 \Rightarrow 6$

Converge  
to base case

# Checking Membership in A List

```
member(_, [ ])      -> false;  
member(H, [H | _]) -> true;  
member(X, [_ | T]) -> member(X, T).
```

What is the arity of member?

What would happen if we swapped the first two lines?

How about the last two?

# Recursively Building Lists

`bump([]) -> [];`

`bump([H | T]) -> [H+1 | bump(T)].`

`bump([2,4,6])  $\Rightarrow$  [ 3 | bump ([4,6]) ]`

`bump([4,6])  $\Rightarrow$  [ 5 | bump([6]) ]`

`bump([6])  $\Rightarrow$  [ 7 | bump([]) ]`

`bump([])  $\Rightarrow$  []`

`[7 | [] ]  $\Rightarrow$  [7]`

`[5 | [7] ]  $\Rightarrow$  [5, 7]`

`[3 | [5, 7] ]  $\Rightarrow$  [3, 5, 7]`

# Some List Functions

```
1> lists:max([1,2,3, 4, 3]).  
4  
2> lists:reverse ([1,2,3]).  
[3,2,1]  
3> lists:split(2, [3,4,10,7,9]).  
{[3,4],[10,7,9]}  
4> lists:zip([1,2,3],[5,6,7]).  
[{1,5},{2,6},{3,7}]  
5> lists:delete(2, [1,2,3,4,2,4,2]).  
[1,3,4,2,4,2]  
6> lists:last([1,2,3,4,2,4,24]).  
24  
7> lists:length([1,2,3,4,2,4,24]).  
** exception error: undefined function lists:length/1  
8> length([1,2,3,4,2,4,24]).  
7
```

length/1 is a BIF



# Homework for Next Week

- Download and install Erlang
- Make sure you understand the results of all examples in the Pop Quiz
  - check them in the shell!
- Open the Erlang Reference Manual
  - downloadable with the code
- Read the man page of the 'lists' module
  - experiment with some list functions
- Solve the problems on the next slide

# Problems for Next Week

1. Write a module 'fib' exporting 'fib/1'  
fib(N) ->  
the Nth element in the Fibonacci series
2. Write mymath:power(x, y) ->  
x to the power of y for integers
3. Write mylists:nth(Nth, List) ->  
the Nth element in List

Not for submission, but I'll ask about them next week 😊

# Next Week

- Sequential Erlang programming
  - control flow – guards, case, if
  - looping with tail recursion
  - list comprehension
  - high order functions, funs